

MSc thesis in Geomatics

Spatio-temporal Transformers for Wildfire Prediction

Carlo Cordes

June 2026

A thesis submitted to the Delft University of Technology in partial fulfillment of the requirements for the degree of Master of Science in Geomatics

Carlo Cordes: *Spatio-temporal Transformers for Wildfire Prediction* (2026)

© This work is licensed under a Creative Commons Attribution 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by/4.0/>.

The work in this thesis was carried out in the:



Geo-Database Management Centre
Delft University of Technology

Supervisors: Dr. Azarakhsh Rafiee
Justin Schembrie

Co-reader: tbd

Contents

1. Related Work	2
1.1. Drivers & Metrics of Wildfire Risk Modeling	2
1.2. The Transformer	3
1.3. Multimodal Learning Tasks in Remote Sensing	7
1.4. Memory Acceleration	8
1.5. Research Gap	8
2. Methodology	10
2.1. Workflow & System Overview	10
2.2. Dataset Construction	10
2.3. Model Architecture	12
2.4. Training Routine	19
2.5. Experiments	22
Bibliography	24

TODOs

TODO 1	Work in progress	6
TODO 2	Work in progress	8
TODO 3	work in progress	10
TODO 4	Update figure of timestamp configurations	10
TODO 5	Complete model architecture diagram and reference	12
TODO 6	Figure of reflective patch padding	12

1. Related Work

1.1. Drivers & Metrics of Wildfire Risk Modeling

Wildfire events are highly complex phenomena. First and foremost, due to its highly violent nature, fires in the wild as well as urban areas pose as one of the most destructive of natural hazards. Furthermore, the stochastic yet extremely rare nature of such events pose as a challenge in quantifying, reasoning and mitigating such events. On top of being difficult to predict, such fire events are highly locally concentrated. In 25 years of recorded data in Spain, a region of 6% of the total landmass is responsible for the majority of outbreaks (Álvarez-Díaz et al., 2015). Having a close pulse on such events is crucial as forest systems serve a crucial importance to establishing a healthy carbon balance and are a fundamental component of mitigation global carbon combustion (Illarionova et al., 2022).

In recent decades, many areas world-wide have been experiencing an increase of such events. While a great deal of the scientific discussion has been centered around correlating land use and risk mitigation to the occurrence of these fires, a large factor can be attributed to a novel set of climate conditions. In the United States, more particularly in the western state of California, a large spike in such events from the 1980s and onwards can be found to correlate significantly with the conditions a warming climate poses (Westerling et al., 2006). With elongated dry summers, drought and stagnating vegetation health, especially regions of mid-elevation are found to be directly be affected by the characteristics of a warming climate.

Assessing and describing the effects of wildfires can be essentially split into two disciplines. Conceptually, a fire disaster describes the propagation of a fire event through combustible material, given the initial start of a fire. The terminology to focus on here are the probability of ignition of a fire and the thereby causal spread of it. Describing these two disciplines separately as well as quantifying their interdependencies lies at the heart of quantifying the risk a fire poses (Castro et al., 2010). In turn, this allows for authorities to make concentrated choices regarding mitigation measures, in which both spark and ignition have to be addressed. In essence, the discipline can be contextualized as the conditional probability of a spread of a fire given an ignition.

Studies find global wildfire ignition with 90% likelihood to be caused by human activity (Arndt et al., 2013). A further investigation of fire ignitions in California, US used 10 years worth of ignition data to show the local dispersion of events caused by humans versus other effects. It showed strong spatial between human-induced events and urban areas as well as remote nature and lightning strikes (Chen & Jin, 2022). It could be shown that this effect is also highly locally optimized and needs to be assessed regionally, once again underlining the importance of good interpretation of local climate health conditions.

Taking ignition probability out of the picture, there are a few crucial factors influencing the extent, intensity and probability of spread. Most importantly, average daily temperatures, correlating with summer season make the biggest effects. Along with a deprecated annual rainfall rate and sparsity thereof, the general distribution of anthropogenic activity is a notable factor. Furthermore, studies have linked gross domestic product (GDP) density to the extent of fire events (Aldersley et al., 2011). Urban areas therefore not only contain more sealed surfaces,

1.1. Drivers & Metrics of Wildfire Risk Modeling

posing as a natural fire barrier but with rising GDP per area, extended mitigation measures could be the underlying factor of less burnt area occurring in such high value producing areas.

Other factors influencing burn extent include natural difference vegetation indices (NDVI). This monitors vegetation health through infrared radiation from satellite stations. Vapor pressure, humidity are contributing factors. A factor that can lead to containment or further acceleration of such events is topography, as terrain is more susceptible to propagating fire in high-slope conditions. Similarly, wind has the analogous ability to shift flame angle to set neighboring combustible material afire. Lastly, determining biomass density and fuel moisture are avid determinants of risk (Rafaqat et al., 2022). Wind is an especially dangerous factor as at high velocities it can also lead to transport of burning materials, sparking new fires elsewhere, a factor that was especially harmful in the January 2025 fires in Los Angeles (Syphard et al., 2022).

1.2. The Transformer

1.2.1. Background & Motivation

Processing data in the machine learning domain, whether text or image, was historically dominated by Recurrent Neural Networks (RNN) or Long-Short-Term Memory (LSTM). These types of learning networks rely on a chain-like structure, processing data in series to make a prediction. Speaking specifically for processing sequences of satellite imagery, a bottleneck arises when scaling to large datasets. As the processing of data sample n_t is dependent on n_{t-1} , such a process is inherently inefficient due to its incompatibility to be parallelized. Furthermore, such models are subject to the *Vanishing Gradient* effect. This describes the loss of information a sequential model encounters as the size of the input sequence increases. Every element is dependent on the one before it, so as a signal passes through larger amounts of computations, information will get lost (Hochreiter, 1998).

A fundamental breakthrough was made in modern technology with the transformer architecture, the goal of which was to eliminate the sequential computation bottleneck. In direct consequence, this would tackle the the issue modeling accuracy being dependent on the distance of elements of the encoded input sequence. The transformer, as originally conceived for nernal translation tasks, was the pioneering idea in this regard. The proposed architecture relies not on processing data sequentially, but treating the input data entirely at once. Inputs no longer rely on their relative position in the sequence, but are much rather position invariant concepts. Every part has the direct ability to interact with all others without relying on those in between (Vaswani, 2017). The regressional analysis was consequently replaced by the attention mechanism, that allows sequence segments to self establish their relevance and corellation to each other.

1.2.2. Architecture

The core idea behind the transformer is to form predictions given an input, depending on the application. As mentioned above, the inital design involved translating text between languages not just word for word, but with strong relational semantic meaning. This same concept was later adapted for a variety of different tasks, perhaps most prominently text generation tasks. Modern chatbots use this very mechanism to from educated prediction of outputs given a user input sequence.

Generative Transformers that produce either images of text from input sequences are pre-trained, meaning their weights are set. Under the hood are various matrix multiplication processes that turn inputs into outputs. Crucial for this to work however is the processing of

input sequences into a structure the transformer architecture can understand. To do this, we consider partitions of the input as so-called *tokens* (Naeve et al., 2024) as shown in Figure 1.1.

Raw: The capital of Italy is
Tokens: [The] [capital] [of] [Italy] [is]

Figure 1.1: Tokenization of an input

Each of the input tokens is then processed via a process called embedding. This simply assigns every token to a high dimensional vector:

$$\text{Italy} \rightarrow \vec{e} = \begin{pmatrix} e_1 \\ e_2 \\ e_3 \\ \vdots \end{pmatrix} \in \mathbb{R}^{d_{\text{embed}}} \quad (1.1)$$

To an untrained model this might appear like a random projection of a word into a vector, however this process is learnable, meaning that the way a token might be transformed into an embedding is optimized as a model trains. As a result, trained models gather their entire understanding of reality by assigning many tokens each a distinct position in a very high dimensional embedding vector space. A trained model not only has a position for every term in its vocabulary, but the relative position of these can directly embody similarities of terms in plain language. For example, consider a pre-trained transformer that holds embeddings for the terms *brother*, *sister*, *aunt* and *uncle* in Figure 1.2.




Figure 1.2: Organization of tokens into vector space and relative positional meanings

The important notion to realize is that a trained transformer is able to relate certain terms to one another. Subtracting the vector of *aunt* and *uncle* will yield a similar vector to the one when applying this logic to *brother* and *sister*, a vector that to the transformer encodes a logic of relative gender. Furthermore, taking the dot product of two vectors will provide information about similarity of these two vectors. Analogously to vector geometry, the dot product of two unit vectors pointing in different directions will yield zero:

$$\hat{e}_i \cdot \hat{e}_j = \delta_{ij} = \{1 \text{ if } i = j, \text{ else } 0 \quad (1.2)$$

Suppose such vectors were describing directions in a coordinate system. A zero dot-product dictates that these two vectors are very dissimilar, as they encode different directions. The

1.2. The Transformer

same applies to the embedding dimension. A large dot product of two embeddings signifies similarity between these terms (Sanderson, 2024a).

1.2.3. Attention

Embedding the tokens of an input sequence like in Figure 1.1, for now only gives a vectoral meaning to each individual one, rather than as a complete input sequence. However, as the core concept here is to make predictions from an input sequence as a whole, a mechanism relating the meaning of each token into a conglomerate understanding is necessary. This exact problem is solved by the concept of attention, a powerful matrix encoder that not only serves as the basis of this project, but is responsible for the immense growth in the artificial intelligence sector (Sanderson, 2024b).

Despite being powerful, this concept is very simple. It determines the relative importance of all embeddings to all other embeddings in the input and subsequently updates each embedding, depending on this relative importance. This step can be conceptualized as tokens absorbing information from others. For example, the embeddings of *green apple* will hold much richer meaning after an attention block as the token *apple* will have **attended** to *green*. Thereafter, the embedding of apple is updated. It now sits in a slightly shifted position in embedding space, carrying this richer, updated information.

Mathematically speaking, the attention block is a scaled matrix multiplication. To apply attention to the input sequence, first let the concatenation operation of all tokens of a sequence into a structure that will be called the Value matrix V :

$$[\text{The}] [\text{capital}] [\text{of}] [\text{Italy}] [\text{is}] \rightarrow \begin{pmatrix} \vdots & \vdots & \vdots & \vdots & \vdots \\ \vec{E}_1 & \vec{E}_2 & \vec{E}_3 & \vec{E}_4 & \vec{E}_5 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix} = V \quad (1.3)$$

For every embedding E_i a query vector (Q_i) and a key vector (K_i) is created via convolution of weighted key and query matrices that are convolved with the embeddings of each token:

$$Q_i = W_Q \cdot \vec{E}_i \quad (1.4)$$

$$K_i = W_K \cdot \vec{E}_i \quad (1.5)$$

Attention is then formulated as scaled the matrix multiplication of these two matrices, using softmax to scale the outputs to the domain $[0, 1]$ and scaled by the value matrix $\vec{V}$ (Vaswani, 2017):

$$\text{Attention}(K, Q, V) = \text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{d_{\text{embed}}}}\right) \cdot V \quad (1.6)$$

This formula acts as a bridge allowing each embedding to inform itself about the contents and relevance of all other embeddings. Every embedding possesses a query Q to which the to other tokens respond with a key value K . The dot product between these two in Equation (1.6) thus defines the relevance of tokens to one another. The greater the output of this dot product, the higher the attention value will be for a token pair. The result of the scaled dot product is multiplied by the value matrix. In simple terms it can be read: one token uses its query vector to get information about another with its key vector. If that information is relevant to it and the dot product is high, a weighted part of the queried token is added onto the embedding of the querying token.

This process is called self attention, as the input sequence is only compared to itself. The analogous process of comparing two sequence, for example in translation tasks. The magnitude to which embeddings attend to one another is entirely governed by the weights in W_k and W_q , which are initialized at random and learned during model training.

The resulting attention values are then used to update the embeddings to hold richer meaning:

$$\vec{E}'_i = \vec{E}_i + A_i \quad (1.7)$$

Attention rarely is carried out in such single operations but in many blocks. The raw embeddings are fed to each of these blocks that might contain many of these operations, individually called heads. Each block and each head is thus able to learn slightly different relational attributes between the embeddings. This is called multi-headed attention (Shaw et al., 2018).

Crucially, the computations of each head and block are independent of each other, making this process highly parallelizable and computationally convenient to compute with large Graphics Processing Units (GPU).

1.2.4. Johnson-Lindenstrauss Lemma & Vector representations

(Johnson & Lindenstrauss, 1984)

1.2.5. Vision Transformers

The concept of using embedded sequences as predictors quickly caught on for vision tasks. The bottleneck resided in establishing an analogy to what a token should represent in images. Using the naive approach in extracting a token for each pixel value in every image quickly scales above computable limits as using N tokens for an image of height h and width w scales quadratically with image resolution ($N = h \times w$). Furthermore, attention processes also scale as $O(n_{\text{tokens}}^2)$ as per Equation (1.6), this would quickly lead to unfeasible computations (Papa et al., 2024). The solution was found in extracting not raw pixel values, but patches of them. By flattening each patch into a vector and projecting this into an embedding dimension an analogous structure to embedding was found (Dosovitskiy et al., 2020). It is hereby important that the embedded values are not the exact pixel values but are the product of a convolution operation. This mechanism outperformed convolutional approaches on predicting labeled such as ImageNet, CIFAR-100 or BTAB.

The transformer notably solved one of the main concerns of convolutional neural network approaches, which involves using a kernel as a sliding window to extract features. This process however, while being an asset at extracting local features, failed to correlate features that might have a large extent or distant relationships in an image (Wang et al., 2025). Because transformers have the ability to self-weight relationships of features, this issue is almost entirely resolved. An investigation on the the large dataset Imagenet-21k could show that vision transformer models consistently outperform convolution approaches, amongst others due to its enhanced feature recognition (Ridnik et al., 2021).

A new challenge with the transformer architecture is that in general, compared to its predecessors, such a model requires an extensive set of pre-labeled data in order not just to learn dependencies of its input, but to generalize well to unseen data (Khan et al., 2023).

Transformers performed well with simple classification tasks, but could initially not perform well in denser tasks, image segmentation being a prime example. This is in part due to the excess of attention calculation a model runs through for one prediction. An adaptation to this was found in the sliding-window (SWIN) which particularly for large-scale images bridged

1.2. The Transformer

the gap of computational time from text tasks to the image interpretation (Liu et al., 2021). By only computing self-attention between non-overlapping local windows rather than the entire set, the approach brought greater efficiency.

An invaluable advancement to vision transformer came with frequent applications to remote sensing data. As satellite data is grid-like it is an ideal application. This coincided with a further enhancement to the pipeline: positional encodings. Here, the set of embedded image patches are pre-weighted, proportionally to their distance to other patches. This serves as a kickstarter to the later cross-patch attention process in which close patches will conduct attention more closely due to their initial similarity induced by the positional encoding (Bazi et al., 2021). Early approaches adapted rotary positional encoding from natural language processing into vision approaches (Heo et al., 2024). Further progress was made with the proposal of using overlaid sinusoidal encodings, that do not just pass absolute encoding but relative ones.

This concept could be further extended the temporal dimension. Input images with multiple steps across time could not only share weights given their relative position, but also would explicitly encode time via the same logic. Images of neighboring timesteps could receive more similar encodings while distant ones could not conduct attention as closely (Arnab et al., 2021; Xu et al., 2019).

1.3. Multimodal Learning Tasks in Remote Sensing

Spatio-temporal attention mechanism find a wide array of applications in the remote sensing sector as with sufficiently large datasets, they tend to outperform other architectures (Wang et al., 2024). Specifically for climate phenomena prediction, learning architectures performed well, even before the initiation of vision transformers. Notably, a bayesian neural network architecture was able to produce numerical predictions of el Niño, taking a multi-modal, multi-year dataset to produce valuable prediction results. The method elegantly predicts sea-surface temperatures under varying lead-time and at different time scales (Luo et al., 2022). The same study was conducted with *Earthformer*, a new transformer architecture that specializes in predicting spatio-temporal datasets (Gao et al., 2022). By dividing datasets not just into patches, but into cuboids via two directional and one temporal dimension, the model uses so-called cuboid-attention to forecast El Niño/Southern Oscillation, outperforming earlier LSTM and convLSTM methods.

1.3.1. Wildfire Applications

With natural catastrophe modeling being heavily reliant on data, machine learning architectures arrived in this sector as well. Notably, a study with the interest of Serbia studied three different deep learning approaches (XGBoost, Kolmogorov-Arnold networks and neural networks) to compare their prediction abilities. By combining a multi-modal dataset to predict historical burn-map, the deep neural network (DNN) could show the best prediction accuracy with 83.4% (Durlević et al., 2025)

As large amounts of resources in the sector are being devoted to wildfire mitigation, large academic interest is focused on predicting wildfire spread (Jain et al., 2020). Specifically for transformers this area seems to be promising. Focusing on fire spread in the North American region, a study could deliver excellent results using a Swin-based transformer mechanism to deliver next-day fire spread estimates (Li & Rad, 2024). Similarly CNN have successfully been combined with encoder-decoder transformer mechanisms to do similar work. By specifically encoding forest fire status into the predicting dataset, better predictive accuracy was achieved. By relying on environmental markers, sociological indicators and fuel conditions, methods like

FiRE-HNL use physics-infused approaches to train models. This discipline could be extended to serve multi-day predictions derived from MODIS & ERA5 data, once again outperforming CNN-approaches with F1 scores of 0.75 (Andrianarivony, 2024; Dubey et al., 2025).

A different sub-discipline focuses on risk metrics with respect to wildfire detection, a factor specifically important for non-urban areas. Another discipline this serves is examining ignition probabilities. By being able to detect wildfires soon after ignition, the root cause of the fire can be reliably identified and used to better understand initial ignition conditions (Dong et al., 2025).

Initial approaches which model ignition and spread probability as a joint risk factor in combination with transformer models look promising. First benchmarks have shown that combining large multi-year datasets of multi-modal satellite imagery can yield promising predicting performance on multiple time scales (Xu et al., 2026).

1.4. Memory Acceleration

1.5. Research Gap

Without a doubt, attention-based transformer is promising for a predictive task. Especially the diverse learnable embedding vector space serves as an asset due to its seemingly infinite application scopes and ability to learn specific tasks. More closely, the storage of natural phenomena as image-based datasets makes the transformer model an ideal application. With the ability of transformer models to learn embedding vector spaces in the natural language processing (NLP) domain (as in Figure 1.2), it begs the question whether these architectures are also able to create analogous representations of climate health, a so-called wildfire-risk embedding space.

First applications to wildfire problems show promising results, however only limited work has been done outside of predicting spread. With some research slowly pivoting towards modeling pure wildfire occurrence risk, first bench-marks have proven the theoretical framework possible with valid results.

Furthermore, little to no information is available regarding discussions of temporal scopes of input and prediction scopes. Short-term multi-day predictions of fire spread are an established discipline however elongating the temporal prediction scope along with modeling pure risk of occurrence have yet to be investigated.

Therefore, the main gap this research addresses, regards temporal flexibility of a predictive mechanism. This includes 3 dependencies: the scope of input data a model uses to form its predictions, the temporal scope of prediction and the dependency on model accuracy to make predictions for the distant future (lead-time). Furthermore, an investigation into feature importance via ablation studies is crucial to understanding predictive performance and will represent a second focus of this work.

1.5. Research Gap

2. Methodology

This section provides an overview of the methodological choices made in the investigation. First, a short account of the system overview is given, before an account for the construction of data into datasets is made. Further, a complete walk-through of the model architecture is given. Finally, the integration of the model into a training workflow is explained.

2.1. Workflow & System Overview

work in progress

```
data:
  spatial_extent:
    longmin: -9.4900
    longmax: -6.0234
    latmin: 36.812
    latmax: 42.2724
  crs: "EPSG:3763"
  scale: 1000 # meters
  temporal_extent:
    start_date: "2020-01-01"
    end_date: "2024-12-31"
    day_interval: 1
  sequence_period: 4 # Time between respective sequences
  sample_extent: 10 # Total time of sample
  sample_period: 1 # Time between images in sample
  target_extent: 4
  lead_time: 0 # Time between last sample and target
```

2.2. Dataset Construction

2.2.1. Timeframe Configuration

A section of the research objective is to assess the influences of the temporal scopes of both data and predictions. Both can be variably adjusted depending on the desired test case. As we handle temporally dynamic datasets we need to define how timesteps of samples are represented in the dataset. To ensure consistency in naming these, we here declare a series of term definitions relevant for this step:

- **Sample:** Input data the model sees to form predictions; describes the set of timestamps for one forward pass
- **Target:** The collection of images, that is variable length, that is used for testing the model predictions
- **Sequence:** A pair of a sample and target
- **Dataset:** A collection of sequences
- **Sequence Period:** Separation between individual unique sequences
- **Extent:** Used both sample and target, defines their respective variable length
- **Lead Time:** The time between the last seen timestamp in the sample and first one in the target. In other words: the amount of time the model predicts into the future

Update figure of timestamp configurations

Every sample is constructed by 3 different kinds of data, that are read, packaged together and channeled into the model:

- **Static:** equivalent in every sample, describes non-variable data (e.g. terrain)




Figure 2.1: Definitions of temporal scope in project datasets

- Dynamic: time-series data consisting of multiple images (e.g. wind)
- Single-dynamic: dynamic data that is less temporally dependent and assumed static over the extent of the sample (e.g. 30-day precipitation sum)

Figure 2.1 illustrates the terms defined above. It shows two consecutive sequences as well as their respective samples and targets. It shows the temporal division of the sample into subsequent timesteps. The distance between sample and target corresponds to the above definition of *lead-time*.

The configuration file that belongs to a model training governs the total time of interest and all parameters defining the temporal relationship of input data. The dataset builder then generates a set of timestamps as governed by the rest of the temporal configurations and iterates through the respective file for every module. Sequences with missing images are skipped entirely to ensure consistency. These dynamic modules are accompanied by the static data, which is the same in every sample. The result is a set of sample/target combinations spanning the entire domain of temporal interest.

2.2.2. Torch Datasets & Tensors

Moving data from images into the Pytorch framework is done using its fundamental data structure: the **Tensor**. Like in the traditional mathematical sense this is an array of arbitrary dimension containing elements of a single data type but has a few computation-driven features. Most importantly, it is optimized to be handled by GPU. Tensors are directly associated with the model instance (elaborated in Section 2.4.2) and have an attribute that defines whether they are adjustable model parameters. This concept is specifically designed for those tensors containing model weights. PyTorch can track the operations imposed on a tensor, which makes later backpropagation to these easier (Pytorch-Contributors, 2019a). The sequences described

in Section 2.2.1 are appended into such a tensor structures and then saved as portable datasets to be made accessible on larger compute resources.

2.3. Model Architecture

In order to motivate the structure of the here applied transformer, this section walks through the sequential steps incoming data takes through the model. This represents one forward propagation. The later showcased training regiment makes this forward pass once for every batch in the training data, thus cycling through it an abundance of times.

2.3.1. Spatial Embedding

The initial objective of the transformer model revolves around turning input data into tokens via patch embedding. The patch size, a model hyperparameter, defines the pixel extent of square patches every input image is divided into. As we handle only single-band data we treat the image as a grid of height and width: $h \times w$.

This is realized via a 2D convolution with stride equal to the desired patch size. This ensures zero overlap between neighbouring patches. An important element to consider is the fact that both height and width of images need not be divisible without rest by the patch size. To counter this we use reflective padding, a process which mirrors input images around their edges to fill in pixel data until the next complete patch. Extending height and width to h_{pad} & w_{pad} allows for perfect image division into patches. As the model input consists of multiple modules M for consecutive time steps t that are processed in batches B of variable sizes the input tensor has dimensionality:

$$\tilde{T} \in \mathbb{R}^{B \times M \times t \times h_{pad} \times w_{pad}} \quad (2.1)$$

Every image of dimension $h_{pad} \times w_{pad}$ is then subject to the flattening operation, here illustrated with a patch size of 16 pixels.

$$P \in \mathbb{R}^{16 \times 16} \rightarrow P_{flat} \in \mathbb{R}^{256} \quad (2.2)$$

Up to this point, pixel values have not changed, simply the indexing is altered. To produce the embedding, the convolution operation then maps every flattened patch into the embedding dimension. Here we showcase the operation for a patch size of 16 pixels, mapping to the embedding of dimension 128:

$$e = wP_{flat} + b \quad (2.3)$$

where

$$w = \begin{pmatrix} w_{1,1} & w_{1,2} & \dots & w_{d_{patch},1} \\ w_{2,1} & w_{2,2} & \dots & w_{d_{patch},2} \\ \vdots & \vdots & \ddots & \vdots \\ w_{d_{embed},1} & w_{d_{embed},2} & \dots & w_{d_{patch},d_{embed}} \end{pmatrix} \in \mathbb{R}^{128 \times 256} \quad (2.4)$$

The convolution weights are learnable and are only shared between samples of the same module. Every module is assigned an independent convolution operation and individual weights.

2.3. Model Architecture

2.3.2. Temporal embedding

The model architecture demands a separate routing of static and dynamic data channels. All dynamic modules with multiple time steps receive a temporal encoding that is an extension to the one described in Section 2.3.1. The objective of this is to give patches that describe the same physical area in different time steps a closer connection to one another. This is done by treating the different time steps for one module together. The single-image 2D convolution is replaced by a 3D convolution that treats the series of images of timesteps t as a 3-dimensional set of values and projects it into the embedding dimension d_{embed} :

$$\mathbb{R}^{t \times d_{patch} \times d_{patch}} \rightarrow \mathbb{R}^{d_{embed}} \quad (2.5)$$

The convolutional operation is analogous to Equation (2.3) for a static embedding with a kernel of dimension:

$$w \in \mathbb{R}^{d_{embed} \times 1 \times P^2} \quad (2.6)$$

In essence, this step results in a token for every patch in every timestep of every image. This effectively patches every timestep in every module independently. The objective here is to prevent the immediate blending of time steps, as this is the objective of the transformer layers at a later stage.

2.3.3. Spatial Embedding

Having obtained a set of embeddings for all images of varying modalities in the model input, this part of the transformer architecture intends to introduce spatial and temporal relationships between the patches called *embeddings*. These greatly influence the magnitude with which patches learn from each other.

The spatial embedding is applied to every patch regardless of being temporally dependency or not. The idea is to assert a learnable idea of proximity from one patch to another. As wildfires are localized phenomena, the the objective of the spatial embedding here is to give neighbouring patches a strong correlation. Distant patches, while sharing lower correlation, can still influence one another, with limited effect. The important difference to a convolution here is that we do not simply rule out distant relationships but keep them contained. Rather than imposing a relationship on the transformer, we give it rough guidelines and freedom to give high weight to the relationship even between distant patches. Having created a notion of spatial similarity is crucial for the later attention process in which more similar patches will conduct higher weighted attention with one another.

In order to create such a relationship we come up with an individual positional encoding for every patch that is added onto each embedding:

$$\vec{e}_{token}' = \vec{e}_{token} + \vec{e}_{pos} \in \mathbb{R}^{d_{embed}} \quad (2.7)$$

To create a positional embedding that retains similarity between proximal patches, we use sinusoidal positional embedding. By imposing sine and cosine waves of varying frequencies over the patch grid, we are able to establish a unique encoding for every patch that has distinct relationships to other patches, depending on their Euclidean distance. This is done by first defining a set of frequencies with $d = \frac{d_{embed}}{4}$:

$$\omega_i = \frac{1}{10000^{\frac{i}{d}}} \quad (2.8)$$




Figure 2.2: Additive encoded values per encoding dimension

The nature of the integer 4 is that in order to create the most unique positional embeddings, we split the added partial embeddings into both height and width and also sine and cosine categories. We scale the frequencies with 10000 in order to fit multiple periods into the image. Interpreting height and width as coordinates we can thus describe the positional encoding of a patch as the vertical concatenation of these frequencies:

$$\vec{e}_{\text{pos}} = \begin{pmatrix} \sin(h\omega) \\ \cos(h\omega) \\ \sin(w\omega) \\ \cos(w\omega) \end{pmatrix} \in \mathbb{R}^{d_{\text{embed}}} \quad (2.9)$$

Figure 2.2 shows samples of the modes of the positional embeddings for every patch in the input image, in this case 80×40 . These include waves that propagate in direction of height and width. For every patch we add the embedding values of all dimensions.

Given the multitude of modes, we can form very individualised encodings for all patches. Figure 2.3 shows the cosine similarity of patches to one another for different embedding dimensions. Here the central patch serves as an anchor-patch, with the values of all patches describing its similarity to it. This shows that sinusoidal positional embedding is able to establish similarities and differences between patches, depending on their respective proximity. Furthermore, it is important to mention that the utility of this process heavily depends on the embedding dimension chosen. A higher embedding dimension allows to pass more sinusoidal modes and hence to create a more distinct characterization of each patch. We here show the cosine similarity for three embedding dimensions, which are more descriptive for larger values.

2.3.4. Temporal Embedding

Weighing patches of the same modality against each other should be entirely up to the transformer model. Afterall it is the idea to find nuanced temporal relationships both within a modularity and between pairs of them. We therefore do not constrain the temporal relationships with weights but much rather provide an additional encoding of time which is added to every token that stems from a dynamic module. Used here is an encoding of the Julian day, which is simply an integer number of the day of the year (0-365). To make use of this in the best way, we encode the Julian day so that it captures the true nature of the climatological conditions: its cyclical nature. The embedding of a patch from a day in December should be close to one in the beginning of January, despite their Julian day being far apart.

2.3. Model Architecture




Figure 2.3: Cosine similarity of patches relative to central anchor-patch at varying embedding dimension

Making such an embedding is realized through a two-step process, first mapping the Julian day J_t into a clock-like cyclical two-dimensional space and then translating it into the embedding dimension. The translation into a cyclical representation v_t is done using a sinusoidal projection via:

$$v_t = \left[\sin\left(\frac{2\pi J_t}{365.25}\right), \cos\left(\frac{2\pi J_t}{365.25}\right) \right] \quad (2.10)$$




Figure 2.4: Julian Day encoding into 2D vector space

Figure 2.4 shows the representation of each Julian day in a two-dimensional vector space, effectively capturing cyclical relationships of months even with distant Julian days. The next step entails moving this to the same embedding dimension which the tokens occupy. Similarly to convolution processes shown earlier, we apply a learnable weight matrix to this vector representation projecting these values into $\mathbb{R}^{d_{\text{embed}}}$:

$$\vec{e}_t = W_{2 \times d_{\text{embed}}} \cdot \vec{v}_t + \vec{b} \quad (2.11)$$

Having embedded all tokens of different modalities into embedding space, we additionally give each token a modality tag $\vec{s}_m$. A tag is unique per module and is a learnable parameter with the objective to give the attention heads a clearer idea of where tokens came from. Throughout the learning process the transformer will be able to differentiate between different modules as the tag injects a numeric meaning to all tokens.

The final transformer-ready embeddings for static and dynamic tokens are additions of their individual spatial, temporal and static embeddings like so:

- Static Embedding: $\vec{z}_n^m = \vec{e}_{\text{token}, n}^m + \vec{e}_{\text{pos}} + \vec{s}_m$
- Temporal Embedding: $\vec{z}_{n,t}^m = \vec{e}_{\text{token}, n}^m + \vec{e}_{\text{pos}} + \vec{e}_t + \vec{s}_m$

Simply adding large vectors to our representations of patches might seem counterintuitive at first, however it is crucial to remember that the representations of data modularities are learnable. The longer the transformer architecture trains, the more it knows to contextualize a patch of precipitation data differently, depending on the time of year. The most elegant analogy to this process is music. If the raw encoding of a patch is analogous to a melody, the positional and temporal encoding represent the harmony accompanying it. Music is propagated as the sums of interfering frequencies, however they are discretely distinguishable to the human ear. The same is true for the transformer, it learns to contextualize the raw encoding (melody) given an extra set of information, the embeddings (harmony).

2.3.5. Token Unification

The heart of the concept in this model relies on early fusion. Rather than separating tokens per module, this approach allows for patches accross time and modularity to freely communicate with all others. Despite having the drawback of computational load here, this is a crucial design choice with the purpose of not overconstraining the system. The token unification step therefore concatenates embeddings into a joint sequence, as a preparation for the next step. Crucially, the order of tokens is of no importance here as the attention steps that follow are invariant to permutation of tokens. Furthermore, due to the spatial and temporal embeddings and embedding tags, tokens are uniquely identifiable. The tensor at this stage holds embeddings accross all timesteps, patches and modules in one place:

$$\vec{T} \in \mathbb{R}^{B \times [h \cdot w(n_{\text{dyn}} \cdot t + n_{\text{static}})] \times d_{\text{embed}}} = \mathbb{R}^{B \times n_{\text{token}} \times d_{\text{embed}}} \quad (2.12)$$

2.3.6. Joint Attention

The tokens in the concatenated tensor at this stage query each others information for the first time. This step can be considered self-attention by mathematical standards, as all embeddings reside in the same tensor, however this really allows embeddings to query one another accross modules. This joint encoder is composed of 8 identical blocks, consisting of 8 attention heads each, equating to a total of 64 attention heads. Each of these have the ability to learn individual feature relationships.

To motivate the design choice of the joint attention block here, a short revisit to the core idea of the original objective of cross-attention for translation learning tasks. A sentence, in which every word forms a token, will likely be rearranged in the translation to another language (see Figure 2.5). In the English language we expect the embedding of *red* to have a large influence that of *fire* and likewise in its French counterpart. When cross-attending these tokens of the two

2.3. Model Architecture

datasets, we expect no consistency in the arrangements in the positions of directly translatable words. Therefore, a cross-attention mechanism considering all tokens accross space and time is mandatory.

EN: The red fire spreads across the land
FR: Le feu rouge se propage à travers le pays

Figure 2.5: Example of token order of two cross-attending datasets

The same concept is true for a spatio-temporal transformer. Rather than separating modularities and conduction isolated self-attention, we consider all possibilities. An example here could be the effect of precipitation on future burn probabilities. When predicting for wildfires, the further the latest rain surge was recorded, the higher the probability for an event will be. This underlines that allowing tokens to attend accross time and modules is crucial at this point.

An indispensable mechanism that is implemented at this position during training is dropout. When learning, models tend to follow their most probable direction of learning, meaning once it locked onto characterizing a specific learnt feature, this is all it will be able to understand. Dropout, systematically works against this, as the goal is not to understand just single relationships between data, but as many as possible. Therefore, at random intervals and attention heads, dropout will temporarily zero-out all attention weights. This signifies to the encoder that for this specific iteration of epoch it will not have access to relating pieces of data in the exact way it is used to. Consequently, it will have to rely on making new connections.

2.3.7. Disaggregation

Having attended the modular tokens, the goal is to reshape the obtained information into a prediction map. Important to consider however is that despite having obtained a set of learned tokens, they still stem from different modalities. In order to make use of dynamic and static data in the best way, the tensor contained all attended global embeddings, is disaggregated back into dynamic and static tensors. The patches from dynamic modules are first averaged accross the temporal dimension:

$$\mathbb{R}^{B \times M \times t \times d_{\text{embed}}} \rightarrow \mathbb{R}^{B \times M \times d_{\text{embed}}} \quad (2.13)$$

It might seem that information is lost via this operation. However, when considering the previous attention block, which has the purpose to weight patches accross timesteps by their importance, this step only aggregates this information into one time step. During training, patches from the more relevant timesteps will receive higher weights, thus the averaged representation of the time series will represent this.

Tokens of static and dynamic layers are consequently concatenated into tensor T_{decoder} , setting up for the decoder:

$$T \in \mathbb{R}^{B \times n_{\text{token}} \times d_{\text{embed}}} \quad (2.14)$$

2.3.8. Decoder Layer

The final task of the vision transformer is to take the embedded vector representations in the embedding space and project them back into a form which represents our ground truth: a 2D image.

It is of course possible to make predictions for every patch. This is however not just an oversimplification but also does not provide the opportunity to predict for as many samples

as possible. Table 1 shows the complete structure of the decoder process that projects back to image space:

Step	Tensor Size $\bar{T}$	Description
Input	$B \times n_{\text{token}} \times d_{\text{embed}}$	Output of attention blocks
Reshape 1	$B \times d_{\text{embed}} \times h_{\text{pad}} \times w_{\text{pad}}$	Reshape back to 4D feature map
Convolution 1 + Batch-Norm / ReLU	$B \times \frac{d_{\text{embed}}}{2} \times h_{\text{pad}} \times w_{\text{pad}}$	3×3 kernel with padding, reduce channel depth
Upsample 1	$B \times \frac{d_{\text{embed}}}{2} \times 4h_{\text{pad}} \times 4w_{\text{pad}}$	$4 \times$ scale, bilinear interpolation
Convolution 2 + Batch-Norm / RELU	$B \times \frac{d_{\text{embed}}}{4} \times 4h_{\text{pad}} \times 4w_{\text{pad}}$	3×3 kernel with padding, reduce channel depth
Upsample 2	$B \times \frac{d_{\text{embed}}}{4} \times 16h_{\text{pad}} \times 16w_{\text{pad}}$	$4 \times$ scale, bilinear interpolation
Convolution 3	$B \times 1 \times 16h_{\text{pad}} \times 16w_{\text{pad}}$	1×1 kernel, pixel-wise classifier
Slice	$B \times 1 \times h \times w$	Un-pad operation to original input size

Table 1: Stepwise decoder mechanism & prediction head

In this stage we use a progressive resolution decoding. After first reshaping the tensor into four dimensions, equal to the dimension of the output we require, a series of convolution and upsample steps are used. Each convolution halves the dimension of the embedding, after which the upsample redistributes this via bilinear interpolation to widen the sizes of the patches back into a full-size image. Finally, the decoder features the *un-pad* operation which discards the padding created initially to prevent data disturbance.

2.3.9. Model Parameters

By design, the total amount of parameters in the model is dependent on the temporal extent of the sample, as every patch and every timestep receives an individual encoding. The model uses 90% of its parameters on encoding while the remainder is equally distributed on embedding and decoding. Table 2 shows the total parameters for a baseline model with a sequence extent of 5 days.

Module Name	Parameter Count
static_embeds	49,280
dynamic_embeds	57,472
join_backbone	1,586,176
sequence_projector	16,512
decoder	92,481
Standalone Parameters	1,536
TOTAL	1,803,457

Table 2: Model Parameters per class

2.4. Training Routine

The setup of the training regimen of models is crucial as without certain guidelines, the model could make faulty predictions. The most important vocabulary in this regard is *overfitting* and *generalization*. Generalization describes the ability of the model to transfer its learned patterns from training data to previously unseen data. This showcases the ability to concretely learn the physical underlying the data, rather than simply memorizing it. A model is said to be overfitting, if the prediction results rely on the training data including all its inaccuracies and noise rather than actual representation of data. The training routine used herein caters to both of these issues. As a guide, Algorithm 1 shows the pseudo-code implementation of the training regiment.

To promote generalization of a trained model, we use a train-validation-test split of 70-15-15 respectively. These are the proportions the dataset is split into at the beginning of training.

```

1 Input: Load dataset, initialize dataloader, move to GPU
2 Calculate class imbalance positive weight  $w$ 
3 Initialize loss function, optimizer, model from config
4 Split data (70/15/15) into training- validation- and testing sets
5  $\text{best\_loss} = \infty$ 
6 for epoch in num_epochs:
7     train_loss = 0
8     for batch in train_data:
9         Predict for batches given static & dynamic inputs
10        Compute losses based on WBCE, predictions & ground truth
11        Set optimizer gradients to zero
12        Backward pass
13        Update weights
14        train_loss += batch_loss
15    Compute: validation loss on validation dataset
16    if validation_loss < best loss:
17        best_loss = validation_loss
18        Save model
19    if no improvement for 7 epochs:
20        break
21    Compute: test loss on test dataset

```

Algorithm 1: Training routine

2.4.1. Loss Function

Following a successful forward pass, we evaluate the accuracy of the model by comparing its prediction to the ground truth as defined by a loss function. The main specification the loss function has to adhere to in this problem set-up is the type of target domain in its binary form (fire / no fire). The other is the lopsidedness of our target categories. Due to the very rare nature of fire events, most of the target data describes no-fire, or to the model: zeros. We therefore choose a categorical loss function that in addition to being sensitive to class imbalances is scaleable by a custom weight, the weighted Binary Cross-Entropy loss function (WBCE). It depends on the model prediction p_i , the ground truth y_i and positive weight w . The loss function sums N comparisons in total, one for every pixel in the ground truth.

$$L = -\frac{1}{N} \sum_{i=1}^N w y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \quad (2.15)$$

This setup allows for customizable weighing of output cases. As a concrete example, Equation (2.15) allows us to treat categorical errors made by the model differently. The penalty given by the loss function should be much higher for an incorrect prediction of high-likelihood of fire than a missed prediction of a fire as the no-fire category dominates in overall frequency. This also prevents the model from simply predicting zero to all images, in which case it would still be correct in most cases, however wildly missing the point of its training exercise.

Figure 2.6 shows the loss plane for all possible prediction/ground-truth combinations and a class imbalance $w = 5$. This illustrates a zero loss for all correct predictions along with the two




Figure 2.6: Loss surface of Weighted Binary Cross Entropy (WBCE) loss function

extreme parts of the domains in which the model has made an error. By predicting less than the actual ground truth the model is consistently penalized harsher than vice-versa.

Parameter choice of the positive weight w is determined by the entirety of the training dataset before training and is then passed as an input parameter. In essence, all ground truth images are summed to yield a balance of categories.

2.4.2. Back Propagation & Gradient Descent

The goal of the training regimen is to produce the combination of model parameters that results in the least amount of loss, more specifically a minimal validation loss. To achieve this, the loss is evaluated against all model parameters after every batch. This optimization problem can be thought of a local or global minimum in the *loss-landscape* of $L(w_1, w_2 \dots w_i)$ to a model parameter combination that produces a lower loss, the following steps work to descend along the steepest gradient of this landscape. To do this, we calculate the gradient of the loss due to each model parameter:

$$g_t = \frac{\partial L}{\partial w_i} \quad (2.16)$$

This tries to approximate the extent to which the loss, which describes the discrepancy between prediction and target, is due to this specific parameter. Parameters with a high gradient shall

be updated more aggressively than those being responsible for little loss. The method for updating these model parameters is accompanied by a further parameter, the learning rate η . This dimensionless value scales the overall aggressiveness of model updates. It's value heavily influences the stability of the training process. If chosen inadequately, the model may never find an equilibrium state. The updated weight parameters θ are sums of their original values and the loss gradient, scaled by the learning rate:

$$\theta_{t+1} = \theta_t + \eta g_t \quad (2.17)$$

2.4.3. Optimizer & Learning Rate Scheduler

Pytorch uses to further implementation to aid this model weight update process. First is the learning rate scheduler, in this case we use the OneCycle learning rate scheduler. This system is designed to dynamically change the learning rate to a dynamic one that changes with every epoch (η_t). The job of this specific scheduler is to start with a lower learning rate and grant the system so-called warm-up epochs. This caters to one specific problem that is common in machine learning model training and specifically one of this kind which deals with high class imbalance. The model weights, as initialized before the first training loop are random and represent no coherence with the data or reality. Applying a full-scale learning rate to these can lead to erratic loss jumps in subsequent epochs. Using the earlier landscape analogy this corresponds to overshooting a downhill section in a landscape. Once the training process has found its momentum, the scheduler will gradually increase the learning rate to a higher value. This is the stage in which the gradient descent process has gained its momentum, and we require computed loss gradients to have maximum effects on the model parameters. It is this exact stage in which the model learns most about its tokens. As the loss function approaches its local minimum the scheduler ideally decreases the learn rate again, as not to overshoot the minimum. The effect of this is a user setting and is proportional to the total amount of epochs the system is in training for.

Another crucial implementation aiding training performance is the optimizer. Here chosen was the *AdamW* optimizer. This creates a valuable synergy with the learning rate scheduler to produce and enhanced image of how weights should be updated, following their loss computation. Along with the gradient a weight produces for the loss, the optimizer tracks momentum vectors $\hat{m}_i$ for each weight over the course of a few epochs. The logic insinuates that a parameter improving in the same trajectory for subsequent epochs, will likely continue in the same direction. It can therefore confidently receive a larger update as compared to a parameter that has recently stagnated or fluctuated. Together with the scheduled learning rate the weight updates postulated in Equation (2.17) becomes:

$$\theta_{i+1} = \theta_i - \eta_t \cdot \left(\frac{\hat{m}}{\sqrt{\hat{v}_t + \varepsilon}} + \lambda \theta_i \right) \quad (2.18)$$

Another parameter $\hat{v}_i$ tracks the erraticism a parameter shows during epochs. Those values behaving unpredictably will receive suppressed updates while those showing consistency will be rewarded greater updates (Kingma, 2014; Pytorch-Contributors, 2019b).

2.5. Experiments

The aim of this study is mainly to create such a spatio-temporal transformer model that serves as a predictor for wildfire events. More specifically, focus will be given to investigating the temporal dependencies of both the input sample and output target. Investigation into how

2.5. Experiments

these drive accuracy will be given. Furthermore, it is crucial to understand feature importance. Therefore, the latter part will focus on extracting feature importance.

To guide the reader through these experiments the following research questions motivate the study of each:

1. **Varying sample/target length:** How does temporal depth influence climatological contextualization and prediction accuracy?
2. **Lead time horizon stress test:** How does prediction vary under extensive lead time (future prediction)?
3. **Ablation studies:** - How does prediction accuracy under constraints of input datasets?

Bibliography

- Aldersley, A., Murray, S. J., & Cornell, S. E. (2011). Global and regional analysis of climate and human drivers of wildfire. *Science of the Total Environment*, 409(18), 3472–3481. <https://doi.org/https://doi.org/10.1016/j.scitotenv.2011.05.032>
- Andrianarivony, H. (2024). *Fire*, 7(12), 482. <https://doi.org/10.3390/fire7120482>
- Arnab, A., Deghani, M., Heigold, G., Sun, C., Lucic, M., & Schmid, C. (2021). ViViT: A Video Vision Transformer. *Corr.* <https://arxiv.org/abs/2103.15691>
- Arndt, N., Vacik, H., Koch, V., Arpaci, A., & Gossow, H. (2013). Modeling human-caused forest fire ignition for assessing forest fire danger in Austria. *Iforest - Biogeosciences and Forestry*, (6), 315–325. <https://doi.org/10.3832/for0936-006>
- Bazi, Y., Bashmal, L., Al Rahhal, M., Dayil, R., & Ajlan, N. (2021). Vision Transformers for Remote Sensing Image Classification. *Remote Sensing*, 13, 516. <https://doi.org/10.3390/rs13030516>
- Castro, F. M. C., Silva, J. S., Fernandes, P., & Rigolot, E. (2010). Towards integrated fire management: outcomes of the European project Fire Paradox. *European Forest Institute*, 23.
- Chen, B., & Jin, Y. (2022). Spatial patterns and drivers for wildfire ignitions in California. *Environmental Research Letters*, 17(5). <https://doi.org/10.1088/1748-9326/ac60da>
- Dong, L., Wang, Y., Li, C., Zhu, W., Yu, H., & Tian, H. (2025). A Joint Transformer–XGBoost Model for Satellite Fire Detection in Yunnan. *Fire*, 8(10). <https://doi.org/10.3390/fire8100376>
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Deghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., & Houlsby, N. (2020). An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. *Corr.* <https://arxiv.org/abs/2010.11929>
- Dubey, P., Rakesh, N., Dubey, P., & Bokoro, P. N. (2025). From Convolution to Attention: Transformer-Based Modeling for Multi-Day Wildfire Spread Forecasting. *IEEE Access*, 13, 128895–128919. <https://doi.org/10.1109/ACCESS.2025.3586922>
- Durlević, U., Ilić, V., & Valjarević, A. (2025). Wildfire Susceptibility Mapping Using Deep Learning and Machine Learning Models Based on Multi-Sensor Satellite Data Fusion: A Case Study of Serbia. *Fire*, 8(10). <https://doi.org/10.3390/fire8100407>
- Gao, Z., Shi, X., Hao Wang, Y. Z., Wang, Y., Li, M., & Yeung, D.-Y. (2022). Earthformer: Exploring Space-Time Transformers for Earth System Forecasting. *Neurips*. <https://doi.org/10.48550/arXiv.2207.05833>
- Heo, B., Park, S., Han, D., & Yun, S. (2024). Rotary Position Embedding for Vision Transformer. <https://arxiv.org/abs/2403.13298>
- Hochreiter, S. (1998). The Vanishing Gradient Problem During Learning Recurrent Neural Nets and Problem Solutions. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, 6(2), 107–116. <https://doi.org/10.1142/S0218488598000094>

- Illarionova, S., Shadrin, D., Tregubova, P., Ignatiev, V., Efimov, A., Oseledets, I., & Burnaev, E. (2022). A Survey of Computer Vision Techniques for Forest Characterization and Carbon Monitoring Tasks. *Remote Sensing*, 14(22). <https://doi.org/10.3390/rs14225861>
- Jain, P., Coogan, S. C., Subramanian, S. G., Crowley, M., Taylor, S., & Flannigan, M. D. (2020). A review of machine learning applications in wildfire science and management. *Environmental Reviews*, 28(4), 478–505. <https://doi.org/10.1139/er-2020-0019>
- Johnson, W. B., & Lindenstrauss, J. (1984). Extensions of Lipschitz mappings into Hilbert space. *Contemporary Mathematics*, 26, 189–206. <https://api.semanticscholar.org/CorpusID:117819162>
- Khan, A., Rauf, Z., Sohail, A., Khan, A. R., Asif, H., Asif, A., & Farooq, U. (2023). A survey of the vision transformers and their CNN-transformer based variants. *Artificial Intelligence Review*, 56(S3), 2917–2970. <https://doi.org/10.1007/s10462-023-10595-0>
- Kingma, D. P. (2014). Adam: A Method for Stochastic Optimization. *Conference Paper at ICLR 2015*. <https://arxiv.org/abs/1412.6980>
- Li, B. S., & Rad, R. (2024). Wildfire Spread Prediction in North America Using Satellite Imagery and Vision Transformer. *2024 IEEE Conference on Artificial Intelligence (CAI)*, 1536–1541. <https://doi.org/10.1109/CAI59869.2024.00278>
- Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., Lin, S., & Guo, B. (2021,). *Swin Transformer: Hierarchical Vision Transformer using Shifted Windows*. <https://arxiv.org/abs/2103.14030>
- Luo, X., Nadiga, B. T., Park, J. H., Ren, Y., Xu, W., & Yoo, S. (2022). A Bayesian Deep Learning Approach to Near-Term Climate Prediction. *Journal of Advances in Modeling Earth Systems*, 14(10), e2022MS003058. <https://doi.org/10.1029/2022MS003058> e2022MS003058 2022MS003058
- Naeve, Z., Mitchell, L., Reed, C., Campbell, P., Morgan, T., & Rogers, V. (2024). Introducing Dynamic Token Embedding Sampling of Large Language Models for Improved Inference Accuracy. *Techrxiv*, 2024(1028), . <https://doi.org/10.36227/techrxiv.173014793.37761346/v1>
- Papa, L., Russo, P., Amerini, I., & Zhou, L. (2024). A Survey on Efficient Vision Transformers: Algorithms, Techniques, and Performance Benchmarking. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12), 7682–7700. <https://doi.org/10.1109/tpami.2024.3392941>
- Pytorch-Contributors. (2019b). *torch.optim*. pytorch. <https://docs.pytorch.org/docs/2.11/optim.html>
- Pytorch-Contributors. (2019a). *torch.tensor*. pytorch. <https://docs.pytorch.wiki/en/tensors.html?>
- Rafaqat, W., Iqbal, M., Kanwal, R., & Song, W. (2022). Study of Driving Factors Using Machine Learning to Determine the Effect of Topography, Climate, and Fuel on Wildfire in Pakistan. *Remote Sensing*, 14, 28. <https://doi.org/10.3390/rs14081918>
- Ridnik, T., Ben-Baruch, E., Noy, A., & Zelnik-Manor, L. (2021,). *ImageNet-21K Pretraining for the Masses*. <https://arxiv.org/abs/2104.10972>
- Sanderson, G. (2024b). *Attention in transformers, step-by-step*. <https://www.3blue1brown.com/lessons/attention>

- Sanderson, G. (2024a). *Transformers, the tech behind LLMs*. <https://www.3blue1brown.com/lessons/gpt/#transformers-the-tech-behind-llms-deep-learning-chapter-5>
- Shaw, P., Uszkoreit, J., & Vaswani, A. (2018). Self-Attention with Relative Position Representations. *Corr.* <http://arxiv.org/abs/1803.02155>
- Syphard, A. D., Keeley, J. E., Gough, M., Lazarz, M., & Rogan, J. (2022). What Makes Wildfires Destructive in California?. *Fire*, 5(5). <https://doi.org/10.3390/fire5050133>
- Vaswani, A. (2017). *Attention Is all You Need*. <https://doi.org/10.48550/arXiv.1706.03762>
- Wang, R., Ma, L., & He, G. (2024). *Transformers for Remote Sensing: A Systematic Review and Analysis*. *Sensors (Basel)*. <https://doi.org/10.3390/s24113495>
- Wang, Y., Deng, Y., Zheng, Y., Chattopadhyay, P., & Wang, L. (2025). Vision Transformers for Image Classification: A Comparative Survey. *Technologies*, 13(1). <https://doi.org/10.3390/technologies13010032>
- Westerling, A. L., Hidalgo, H. G., Cayan, D. R., & Swetnam, T. W. (2006). Warming and Earlier Spring Increase Western U.S. Forest Wildfire Activity. *Science*, 313(5789), 940–943. <https://doi.org/10.1126/science.1128834>
- Xu, D., Ruan, C., Kumar, S., Körpeoglu, E., & Achan, K. (2019). Self-attention with Functional Time Representation Learning. *Corr.* <http://arxiv.org/abs/1911.12864>
- Xu, Z., Cheng, S., Wang, L., He, H., Sun, W., Li, J., & Xu, L. L. (2026). BCWildfire: A Long-term Multi-factor Dataset and Deep Learning Benchmark for Boreal Wildfire Risk Prediction. *Proceedings of the AAAI Conference on Artificial Intelligence*, 40(46), 39486–39494. <https://doi.org/10.1609/aaai.v40i46.41299>
- Álvarez-Díaz, M., González-Gómez, M., & Otero-Giraldez, M. S. (2015). Detecting the socioeconomic driving forces of the fire catastrophe in NW Spain. *European Journal of Forest Research*, 123, 1087–1094. <https://doi.org/10.1007/s10342-015-0911-1>

